

gibbs_sampler

January 22, 2023

0.0.1 Imports

```
[1]: import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
from tqdm import tqdm
from typing import Tuple, Callable, Optional, TypedDict
```

0.0.2 Main part

Task 1 Calculations needed for task 1 and 9 can be found at the end of pdf with report written out by hand.

Task 2

```
[2]: def gibbs_sampler(num_samples: int, previous: Tuple[int, int] = (0, 0), thin_out:
    ↳ int = 1, burn_in: int = 0) -> pd.DataFrame:
    ret = [previous] # Since these are Markov chains we only need to remember
    ↳ the last generated item
    for t in range(burn_in + num_samples * thin_out):
        idx = int(np.random.rand() > .5)
        cloud_prob = 4./9. if previous[0] else 1./21.
        rain_prob = .815 if previous[1] else .216
        previous = (np.random.choice((0, 1), p=[1 - rain_prob, rain_prob]),
    ↳ previous[1]) if idx == 0 else (previous[0], np.random.choice((0, 1), p=[1 -
    ↳ cloud_prob, cloud_prob]))
        if t > burn_in:
            if (t - burn_in) % thin_out == 0:
                ret.append(previous)

    return pd.DataFrame(ret, columns=["Rain", "Cloudy"])
```

```
[3]: N = 100
gibbs_sampler(N)
```

```
[3]:      Rain  Cloudy
0         0         0
1         0         0
2         0         0
```

3	0	1
4	1	1
..
95	1	1
96	1	0
97	0	0
98	0	0
99	0	0

[100 rows x 2 columns]

Task 3

```
[4]: probs = gibbs_sampler(N).mean()
      print(f"P(R = T | S = T, W = T) is approximately: {probs['Rain']}")
```


P(R = T | S = T, W = T) is approximately: 0.21


Task 4 & 5

```
[5]: def freq_by_iter(num_iter: int, **kwargs) -> pd.DataFrame:
      gibbs = gibbs_sampler(num_iter, **kwargs)
      return gibbs.cumsum().divide(gibbs["Rain"].index.values, axis=0)

N = 50000
print(f"P(R = T | S = T, W = T) approximated by sampling 50 000 samples:␣
      ↳{gibbs_sampler(N).mean()['Rain']}")
for _ in range(2):
    freq_by_iter(N).plot(title="Frequency by iteration", ylabel="Frequency",␣
      ↳xlabel="Iteration", figsize=(10, 10))
    plt.show()
```

P(R = T | S = T, W = T) approximated by sampling 50 000 samples: 0.32046





Based on the plot above I would suggest burn-in time of at least 7500 iterations. However, if we want to be sure the burn-in time is sufficient 10000 iterations may be better idea.


Task 6

```
[6]: def autocorr_plot(samples: pd.DataFrame, max_lag: int = 100):
    autocorrs = [(samples["Rain"].autocorr(i), samples["Cloudy"].autocorr(i))]
    for i in range(1, max_lag + 1):

    plt.figure(figsize=(10, 10))
    plt.plot(autocorrs, label=["Rain", "Cloudy"])
    plt.title("Autocorrelation by size of window")
    plt.xlabel("Window size")
    plt.ylabel("$\\rho$")
```

```
plt.legend()
plt.show()

autocorr_plot(gibbs_sampler(N), 100)
```



Based on this plot I would suggest interval of 20 iterations for drawing approximately independent samples.

Task 7 See implementaion of `gibbs_sampler` above (Task 2).

Task 8

[7]:

```

print(f"P(R = T | S = T, W = T) approximated with burn-in and thinning-out:␣
↪{gibbs_sampler(100, burn_in=10000, thin_out=20).mean()['Rain']},␣
↪\napproximation in Task 3 was: {probs['Rain']}")

```

P(R = T | S = T, W = T) approximated with burn-in and thinning-out: 0.37,
approximation in Task 3 was: 0.21


Due to small number of samples both approximations are quite random, however approximation with burn-in and thinning-out is more likely to be close to the real value of this probability which is 0.32. See histograms below:

```

[8]: vanila = []
      burn_thin = []
      large = []
      large_burn = []
      for _ in range(100):
          vanila.append(gibbs_sampler(100).mean()["Rain"])
          burn_thin.append(gibbs_sampler(100, burn_in=10000, thin_out=20).
↪mean()["Rain"])
          large.append(gibbs_sampler(50000).mean()["Rain"])
          large_burn.append(gibbs_sampler(50000).mean()["Rain"])

      plt.figure(figsize=(10, 10))
      plt.hist(burn_thin, label="with burn-in and thin-out", alpha=.5)
      plt.hist(vanila, label="w/o burn-in or thin-out", alpha=.5)
      plt.hist(large, label='w/o burn-in or thin-out, large number of samples', alpha=.
↪3)
      plt.hist(large_burn, label='with burn-in or thin-out, large number of samples',␣
↪alpha=.3)
      plt.ylabel("Number of approximations")
      plt.xlabel("Approximated value")
      plt.title("Histograms of approximations")
      plt.legend()
      plt.show()

```



In those we can see that approximation by sampling with burn-in and thin-out is concentrated around real value 0.32 while the one without burn-in and thin-out is much more diverse. As we can see taking larger number of samples gives better approximation even without burn-in or thinning out (and even better with them).

Task 10 - Convergence diagnostics

```
[9]: def gelman_rubin_samples(samples: np.array) -> np.array:
      """
      Calculates Gelman-Rubin statistic for chain from which came samples.
      :param samples: array of shape (m, n, k) where m - number of chains, n-
      ↪ number of samples, k- number of batched chains in sample (i.e. 2 when we work
      ↪ with sampler from SAD project)
```

*:return: array of length k containing Gelman-Rubin statistics for each of
↳ batched chains.*

```

"""
m, n, k = samples.shape
W = np.empty(k)
B = np.empty(k)
sigma_hat = np.empty(k)
for i in range(k):
    sample = samples[:, :, i]
    W[i] = np.var(sample, axis=1).mean()
    B[i] = np.var(sample.mean(axis=1))
    sigma_hat[i] = (n-1)/n * W[i] + B[i]
V_hat = sigma_hat + B/(n * m)
return np.sqrt(V_hat/W)

```

```

def gelman_rubin(sampler: Callable[[int, Optional[Tuple[int, int]],  

↳ Optional[int], Optional[int]], pd.DataFrame], num_iter: int, num_samples: int,  

↳ = 3, **kwargs: TypedDict) -> float:
    samples = np.stack([sampler(num_iter, **kwargs).to_numpy() for _ in  

↳ range(num_samples)])
    return gelman_rubin_samples(samples)

```

[10]: `def gelman_plot(samples: np.array):`

*"""
Generates plot of Gelman-Rubin statistic. Function needs chains of length
↳ greater than 50.*

*:param samples: Samples from m markov chains
"""*

```

assert samples.shape[1] > 50, "More than 50 samples needed from each chain"
ret = []
for i in tqdm(range(40, samples.shape[1])):
    ret.append(gelman_rubin_samples(samples[:, :i, :]))

```

```

plt.figure(figsize=(10, 10))
plt.plot(ret, label=["Rain", "Cloudy"])
plt.ylabel("Shrink factor")
plt.xlabel("Iteration")
plt.title("Gelman-rubin statistic")
plt.legend()
plt.ylim(bottom=1., top=1.01)
plt.show()

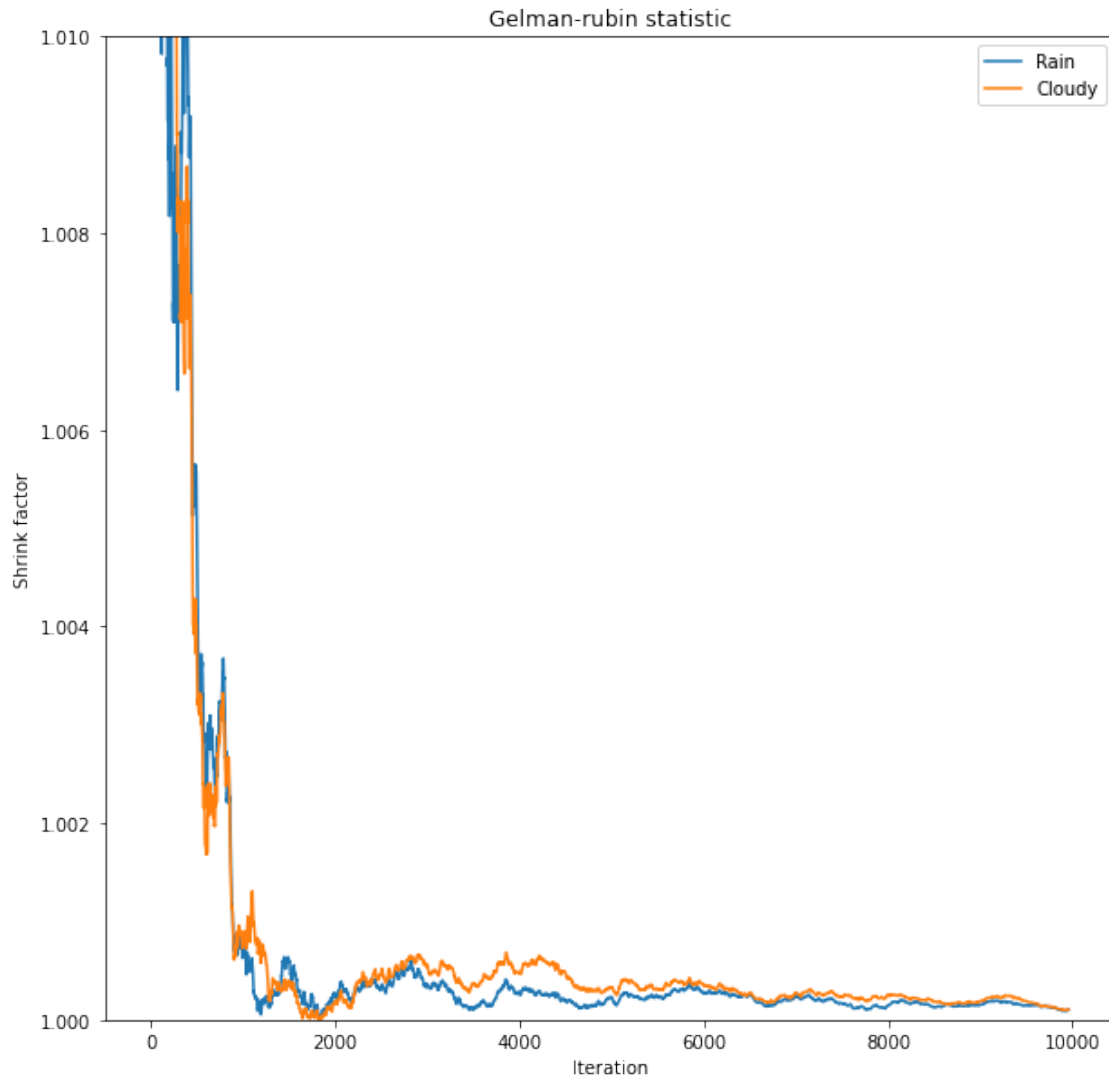
```



```
[11]: samples = 10000
sample = np.stack([gibbs_sampler(samples, previous=(1, 1)).to_numpy() for _ in
↳ range(7)])

gelman_plot(sample)
```

100%| 9960/9960 [00:04<00:00, 2246.47it/s]



In this plot we can see that with number of iterations going up value of statistic is converging to 1, so we can assume that the chains are convergent.